

Saketh Vadlamudi

Hyderabad | 9704808576 | sakethvadlamudi22@gmail.com | [LinkedIn](#) | [GitHub](#) | [Website](#)

Summary

AI/ML Engineer with over 4 years of experience building production-grade GenAI, RAG, and ML systems at scale. Designed multi-agent reasoning engines, hybrid RAG search systems, and deployed cloud-native ML pipelines. Delivered measurable impact: 60%+ latency cuts, 2× performance gains, and 65+ engineer hours saved monthly. Focused on designing robust, scalable AI systems driving product outcomes.

Technical Skills

- **Programming languages:** Python, SQL
- **AI/ML:** Generative AI, Machine Learning (ML), Natural Language Processing (NLP), Retrieval Augmented Generation (RAG), Transformers, Agentic AI, LangGraph, LangChain, Pydantic, Fine-tuning (QLoRA)
- **Cloud Technologies:** AWS (lambda, DynamoDB, Redshift, SageMaker, S3, EC2, ECR), GCP (BigQuery, Vertex AI)
- **Tools:** Git, GitHub Actions, Docker, FastAPI, Unix, Django, Redis, Jenkins, Kubernetes

Professional Experience

Oracle

May 2026 – Present (Hyderabad)

Senior Data Scientist (AI Applications Engineer)

- **Architected a scalable enterprise NL-to-SQL** platform using **Codex Skills**, converting natural-language questions into secure Oracle SQL and business insights; earned **strong leadership approval**.
- **Built a data-centre control dashboard** integrating capacity, supply, server, and network data to support infrastructure planning.
- **Tech Stack** : Python, OpenAI Codex, Docker, Jenkins, Kubernetes, CI/CD

Tiger Analytics

May 2025 – Mar 2026 (Hyderabad)

AIML Senior Associate

- Architected an autonomous, **multi-agent** reasoning engine using **LangGraph** to decompose ambiguous user requests into parallel, dependency-aware execution graphs, **reducing reporting latency by over 60%**.
- Engineered the engine as a **self-correcting, schema-aware** pipeline handling **2000 daily queries** where orchestrated agents autonomously produce **natural language insights, structured data, and interactive Charts** from a single user query.
- **Tech Stack** : Python, LangGraph, Pydantic, AWS (Lambda, DynamoDB, RedShift, Bedrock), FastAPI, PostgreSQL

ValueLabs

Mar 2022 – May 2025 (Hyderabad)

Senior Software Engineer (Data Scientist)

- Developed a Machine Learning (ML) model for **Sentiment Analysis**, which accurately classifies customer reviews on products. Achieved 93% accuracy; insights were adopted by product stakeholders to refine customer messaging.
- Architected an **Advertisement Recommendation System** using collaborative filtering, achieving **MAE of 0.79** and improving ad relevance, leading to a 20% increase in click-through rates.
- Spearheaded **Python processes on Factory Design Mechanism**, reduced run time performance by **2x and 12%** decrease in Campaign drop off rate.
- Implemented several **Python-based automations** for handling manual/regular tasks resulted in significant **reductions** of over **65 working hours** monthly with notable recognition from stakeholders.
- Delivered **4** high-quality, scalable, and secure **Web applications using Django** framework by creating innovative, user-centric solutions that have consistently exceeded client expectations.
- **Tech Stack** : Python, ML, NLP, Django, Unix, AWS (EC2, ECR, Sagemaker, S3), BigQuery, Docker, Jira

Additional Projects

Generative AI :- HybridSeek ([Link](#)) – Semantic RAG Search Engine

- Built a hybrid RAG pipeline (BM25 + dense vector + RRF) for searching over 10k+ document chunks.
- Designed a section-aware chunking + 1024-dim embedding pipeline, achieving >90% semantic recall@10 during evaluation.
- Reduced repeated-query latency by ~8× using Redis caching with exact-match hashing.
- Automated ingestion and re-indexing with Airflow DAGs, processing 1k+ new documents/day.
- Integrated streaming LLM inference with Ollama, delivering sub-3s time-to-first-token latency.

Optimisation :- FineTuning_LLAMA_3.2_3B on Domain Specific Data ([Link](#))

- Architected a fine-tuning pipeline to benchmark proprietary models (GPT-4o) against open-source alternatives (Llama-3.2-3B) for a high-volume e-commerce pricing task.
- Fine-tuned Llama-3.2 with QLoRA + Unsloth on a single NVIDIA T4, cutting memory by 60% and doubling training speed.
- The fine-tuned 3B model achieved a Mean Absolute Error (MAE) of \$47, significantly beating the GPT-4o zero-shot baseline of \$76 and the human baseline of \$127.

Education

Executive Post Graduation, Data Science and AI , IIT Roorkee

2024 - 2025

Bachelor of Technology, Computer Science, Narasaraopeta Engineering College (CGPA 8.94)

2018 – 2022

Awards & Certifications

- ValueLabs CEO Choice Award Winner 2024 ([Link](#))
- Google Cloud Certified Professional Machine Learning Engineer ([Link](#))